

06
PHASE

Data Engineering Fundamentals

Build Reliable Pipelines.
Move Data. Power Insights.



Data
Integration



Data
Transformation



Data Quality
& Reliability



Scalability &
Performance



Security &
Best Practices



Build Strong Pipelines. Enable Better Decisions.

Table of Contents

8 Units • 41 Topics • A Complete Data Engineering Journey

Unit 1 Data Integration (ELT & ETL) 09 – 24

- 1.1 ETL vs. ELT: Why modern cloud stacks prefer Extract-Load-Transform
- 1.2 Batch vs. Real-time processing: When to use which
- 1.3 Change Data Capture (CDC): How databases sync only new/updated rows
- 1.4 Managed Connectors: Using Fivetran, Airbyte, or Stitch to automate ingestion
- 1.5 API Ingestion: Understanding how data is pulled from SaaS tools (Salesforce, Zendesk)
- 1.6 File-based Ingestion: Handling CSV, Parquet, and Avro files

Unit 2 Storage Architecture (Lakes, Warehouses, & Lakehouses).... 25 – 43

- 2.1 Data Warehouse: Structured storage for BI (BigQuery, Snowflake)
- 2.2 Data Lake: Unstructured storage for raw data (AWS S3, Google Cloud Storage)
- 2.3 The Data Lakehouse: The “best of both worlds” (Databricks, Iceberg)
- 2.4 Object Storage vs. Block Storage: Understanding S3/GCS buckets
- 2.5 Storage Formats: Row-based (CSV) vs. Columnar (Parquet) performance

Unit 3 Cloud Data Warehousing (The Big Four) 44 – 66

- 3.1 Google BigQuery: Serverless architecture and “Querying the Web”
- 3.2 Snowflake: Separation of Storage and Compute; Data Sharing features
- 3.3 AWS S3 + Athena: “Query-on-the-fly” logic for data lakes
- 3.4 Azure Synapse: Integrating Power BI with Enterprise DW
- 3.5 Cost Management: Understanding “Slot” usage, “Compute Credits,” and “Scan costs”

Unit 4 Analytics Engineering with dbt (Data Build Tool) 67 – 93

- 4.1 The dbt Workflow: Developing, Testing, and Documenting
- 4.2 Models & Modular SQL: Breaking complex queries into smaller, reusable pieces
- 4.3 Materializations: Views vs. Tables vs. Incremental models
- 4.4 Version Control: Basics of Git (Commit, Push, Pull Request) for analysts
- 4.5 Jinja & Macros: Using Python-like logic inside SQL to automate code
- 4.6 Lineage: Understanding how data flows from Raw to Product tables

Unit 5 Orchestration & Automation (Airflow)94 – 119

- 5.1 The DAG (Directed Acyclic Graph): Understanding pipeline dependencies
- 5.2 Scheduling: Cron jobs and interval-based triggers
- 5.3 Task Failures & Retries: How pipelines recover from errors

- 5.4 Operators & Sensors: Logic for waiting for a file to arrive before running SQL
- 5.5 Airflow vs. dbt Cloud: Choosing the right tool for automation

Unit 6 Schema Design & Data Modeling 120 – 145

- 6.1 Star Schema: Fact tables and Dimension tables
- 6.2 Slowly Changing Dimensions (SCD)
- 6.3 Denormalization: Why cloud warehouses prefer wide tables over complex joins
- 6.4 Surrogate Keys: Generating unique IDs for warehouse records
- 6.5 Partitioning & Clustering: How to organize tables to save money on queries

Unit 7 Data Quality, Reliability, & Observability 146 – 172

- 7.1 Data Validation: Using dbt tests for Uniqueness and Not-Null constraints
- 7.2 Freshness Alerts: Detecting when data hasn't been updated in 24 hours
- 7.3 Schema Evolution: What happens when a source column name changes?
- 7.4 Anomalies: Detecting sudden spikes or drops in record counts
- 7.5 Data Documentation: Maintaining a Data Dictionary for the business

Unit 8 Advanced Data Patterns (Streaming & Privacy) 173 – 199

- 8.1 Streaming Basics: Kafka and Pub/Sub (Real-time events)
- 8.2 Lambda vs. Kappa Architecture: Handling both historical and live data
- 8.3 Data Privacy (GDPR/CCPA): Masking PII (Personally Identifiable Information)
- 8.4 Data Access Control: Administering roles and permissions in the warehouse

Each unit includes a Quick Reference Cheat Sheet at the end

Learnlytics

.handbook

LEARN SMART. ANALYZE BETTER. GROW FASTER.