

03
PHASE

SQL & Databases

Query Data. Manage Databases.
Drive Insights.



Design
Databases



Write
Queries



Filter &
Aggregate



Analyze
Data



Ensure Data
Integrity



Structured Data. Smarter Decisions.

Table of Contents

10 Units • 100 Topics • A Complete Analytics Journey

Unit 1 The Core Foundation (SELECT, WHERE, ORDER BY) 09 – 24

- 1.1 Basic SELECT syntax: columns & aliases
- 1.2 SELECT DISTINCT to remove duplicates
- 1.3 Filtering rows with WHERE clause
- 1.4 Comparison operators: =, <>, >, <, >=, <=
- 1.5 Logical operators: AND, OR, NOT
- 1.6 BETWEEN, IN, and NOT IN filters
- 1.7 Wildcard filtering with LIKE (% , _)
- 1.8 Sorting results with ORDER BY (ASC / DESC)
- 1.9 Limiting results: LIMIT, TOP, FETCH NEXT
- 1.10 Column aliasing with AS

Unit 2 Reality Check (NULL Handling & Type Casting) 25 – 42

- 2.1 What NULL means: unknown vs empty
- 2.2 IS NULL and IS NOT NULL filters
- 2.3 COALESCE to replace NULLs with defaults
- 2.4 NULLIF to convert a value into NULL
- 2.5 CAST and CONVERT for explicit type conversion
- 2.6 TRY_CAST and TRY_CONVERT for safe conversions
- 2.7 Implicit vs explicit type casting pitfalls
- 2.8 ISNULL and NVL: dialect-specific alternatives
- 2.9 NULLs in math: what happens when you add 5 to NULL?
- 2.10 Handling NULLs in logical conditions

Unit 3 Summarizing Data (GROUP BY & Aggregations) 43 – 60

- 3.1 Aggregate functions: SUM, COUNT, AVG, MIN, MAX
- 3.2 GROUP BY syntax and how it works
- 3.3 Filtering grouped results with HAVING
- 3.4 WHERE vs HAVING: key differences
- 3.5 COUNT(*) vs COUNT(column) vs COUNT(DISTINCT)
- 3.6 Grouping by multiple columns
- 3.7 ROLLUP for subtotals & grand totals
- 3.8 CUBE for multi-dimensional aggregations
- 3.9 GROUPING SETS for custom aggregation levels
- 3.10 NULLs in aggregate functions: what gets ignored?

Unit 4 Combining Data Horizontally (JOINS) 61 – 80

- 4.1** Understanding relational JOIN logic
- 4.2** INNER JOIN: matching rows only
- 4.3** LEFT JOIN: all left rows + matched right rows
- 4.4** RIGHT JOIN: all right rows + matched left rows
- 4.5** FULL OUTER JOIN: all rows from both tables
- 4.6** CROSS JOIN: cartesian product use cases
- 4.7** SELF JOIN: joining a table to itself
- 4.8** Joining on multiple conditions
- 4.9** Joining more than two tables
- 4.10** Common JOIN mistakes & NULLs in results

Unit 5 Combining Data Vertically (Set Operations)81 – 102

- 5.1** What are set operations and when to use them
- 5.2** UNION to combine results from two queries
- 5.3** UNION vs UNION ALL: duplicates & performance
- 5.4** INTERSECT to find common rows between queries
- 5.5** EXCEPT / MINUS to find rows in one but not another
- 5.6** Column count & data type rules for set operations
- 5.7** Ordering results from set operations
- 5.8** Using set operations vs JOINS: when each fits
- 5.9** Stacking multiple result sets for reporting
- 5.10** Real-world use case: comparing daily snapshots

Unit 6 Advanced Logic (Subqueries & CTEs) 103 – 124

- 6.1** What is a subquery and where to use it
- 6.2** Subqueries in SELECT (scalar subqueries)
- 6.3** Subqueries in WHERE (filtering with results)
- 6.4** Subqueries in FROM (derived tables)
- 6.5** Correlated vs non-correlated subqueries
- 6.6** EXISTS and NOT EXISTS operators
- 6.7** CTE syntax with the WITH clause
- 6.8** Multiple CTEs in a single query
- 6.9** Recursive CTEs for hierarchical data
- 6.10** CTEs vs subqueries: readability & performance

Unit 7 Professional Analysis (Window Functions) 125 – 147

- 7.1** What are window functions and OVER clause
- 7.2** PARTITION BY to define window groups
- 7.3** ORDER BY inside OVER for row ordering
- 7.4** ROW_NUMBER for unique sequential ranking
- 7.5** RANK and DENSE_RANK: handling ties
- 7.6** NTILE for bucketing rows into groups

- 7.7 LAG to access previous row values
- 7.8 LEAD to access next row values
- 7.9 FIRST_VALUE and LAST_VALUE in a window
- 7.10 Running totals & moving averages (ROWS BETWEEN)

Unit 8 Custom Logic & Cleaning (CASE, Dates, Strings) 148 – 170

- 8.1 CASE WHEN syntax (Simple vs Searched)
- 8.2 CASE WHEN inside aggregates (Pivot logic)
- 8.3 Date data types and current date functions
- 8.4 EXTRACT, DATEDIFF, and DATE_TRUNC
- 8.5 String manipulation: CONCAT, SUBSTRING, TRIM
- 8.6 String searching: CHARINDEX, LIKE, REGEX
- 8.7 Splitting and replacing strings
- 8.8 Building dynamic labels and categories
- 8.9 Handling time zones in analytics
- 8.10 Formatting data for final reports

Unit 9 Performance & Scale (Optimization & Cloud) 171 – 193

- 9.1 How indexes work (Clustered vs Non-clustered)
- 9.2 Reading execution plans (EXPLAIN ANALYZE)
- 9.3 Identifying full table scans vs index seeks
- 9.4 SARGable vs non-SARGable conditions
- 9.5 BigQuery & Redshift architecture
- 9.6 Partitioned & clustered tables in the cloud
- 9.7 Cost management: bytes billed & dry runs
- 9.8 Distribution keys and sort keys
- 9.9 Loading data: GCS/S3 to Cloud SQL
- 9.10 Scheduled queries & materialized views

Unit 10 Building Systems (Procedures, Views & Design) 194 – 220

- 10.1 Creating, replacing & dropping Views
- 10.2 Materialized views and refresh strategies
- 10.3 Stored procedure syntax and structure
- 10.4 Input/Output parameters and variables
- 10.5 Transactions: COMMIT and ROLLBACK
- 10.6 Normalization (1NF, 2NF, 3NF)
- 10.7 Denormalization for analytics performance
- 10.8 Star schema: fact & dimension tables
- 10.9 Entity-Relationship (ER) diagrams
- 10.10 Choosing data types for storage efficiency

Each unit includes a Quick Reference Cheat Sheet at the end